

Linear Regression

(A) Rewrite the WLS objective functions in terms of vectors and matrices as follows:

$$\hat{\beta} = \arg \min_{\beta \in \mathbb{R}^p} \sum_{i=1}^N \frac{w_i}{2} (y_i - x_i^T \beta)^2 = \arg \min_{\beta \in \mathbb{R}^p} \frac{1}{2} (y - X\beta)^T W (y - X\beta) \quad (1)$$

Since W is a diagonal matrix of weights (so symmetric), then

$$\begin{aligned} & (y - X\beta)^T W (y - X\beta) \\ &= (y^T - \beta^T X^T) W (y - X\beta) \\ &= y^T W y - y^T X \beta - \beta^T X^T W y + \beta^T X^T W X \beta \\ &= y^T W y - (X^T W y)^T \beta - \beta^T X^T W y + (X\beta)^T W X \beta \end{aligned}$$

To get the minimum of the above equation, take derivative of β with the following formulas

$$\begin{aligned} \frac{\partial x^T a}{\partial x} &= \frac{\partial a^T x}{\partial x} = a, \quad \frac{\partial x^T b x}{\partial x} = (b + b^T)x \\ \Rightarrow \frac{\partial [(y - X\beta)^T W (y - X\beta)]}{\partial \beta} &= -2X^T W y + 2X^T W X \beta = 0 \\ \Rightarrow (X^T W X) \hat{\beta} &= X^T W y \end{aligned}$$

(B) Numerically speaking, I do not think inversion method is the fastest and most stable way to solve the linear system. Computing and applying the inverse matrix or pseudoinverse is a tremendously bad idea, since it is much more expensive and often numerically less stable than applying other algorithms, especially for high dimensional matrices.

Here are several ways to solve a system of linear equations $Ax = b$ which provide more stability and are computationally efficient compared to inversion method.

- (a) Matrix decomposition: Gaussian or Gauss-Jordan elimination (considered as LU decomposition), Cholesky decomposition, QR decomposition, or SVD, and
- (b) Iterative method: conjugate gradient method.

Which method is optimal depends on the size and properties of the system matrix.

- (a) LU requires A to be square and performs well when A is sparse; it can be used for most linear systems.
- (b) Cholesky performs well for **Hermitian positive-definite matrix** A ($A = LL^*$, where L is a lower triangular matrix with real and positive diagonal entries, and L^* denotes the conjugate transpose of L)
- (c) QR decomposition requires A has linearly independent columns.
- (d) Conjugate gradient requires A to be symmetric and positive definite; it performs well when

A is sparse and too large to be inverted directly for Cholesky.

(C) Since $X^T W X = X^T (W^{1/2})^T W^{1/2} X \stackrel{\text{set } W^{1/2}=V}{=} (VX)^T VX$ is positive-definite, my method will be

the Cholesky decomposition. To solve $(X^T W X) \hat{\beta} = X^T W y$, here is the pseudo-code:

Function inputs:

X: N x P matrix

Y: N x 1 vector of responses

W: N x N diagonal matrix of weights

Function outputs:

$\hat{\beta}$: P x 1 vector of coefficient estimates

Pseudo-code:

- 1) Set $A = X^T W X$;
- 2) Set $b = X^T W y$;
- 3) Set U = Cholesky decomposition of A . Only the upper triangular part of A is used in R so that $A = U^T U$ when A is symmetric.
- 4) Solve $U^T z = b \Rightarrow z = (U^T)^{-1} b$;
- 5) Solve $Ux = z \Rightarrow x = U^{-1} z$;
- 6) Return $\hat{\beta}$ as $\hat{\beta} = x$.

From the following R output, the inverse method is fastest for very small N and P values, but as N and P increase, LU and Cholesky methods perform much more quickly than inverse. LU is the most efficient method of the three and it is a partial Gaussian elimination method.

```
> res # Display benchmarking results
$`N=20, P=5`
Unit: milliseconds
      expr      min       lq      mean     median        uq      max neval
inv_method(X, w, y) 0.075834 0.080395 0.1765681 0.084102 0.086668  4.808907   100
lu_method(X, w, y)  0.161932 0.170484 0.3659880 0.179037 0.194717 17.948158   100
chol_method(X, w, y) 0.160221 0.165923 0.2938887 0.175616 0.194717  9.791719   100

$`N=100, P=25`
Unit: milliseconds
      expr      min       lq      mean     median        uq      max neval
inv_method(X, w, y) 0.629480 0.662265 0.8013725 0.6987570 0.7557745  2.577788   100
lu_method(X, w, y)  0.274827 0.287942 0.3764907 0.3312755 0.3646315  2.410726   100
chol_method(X, w, y) 0.321582 0.346670 0.4112489 0.3703330 0.4079645  0.713297   100

$`N=600, P=150`
Unit: milliseconds
      expr      min       lq      mean     median        uq      max neval
inv_method(X, w, y) 98.66754 100.70594 103.65010 101.69606 102.99835 228.0673   100
lu_method(X, w, y)  13.00013 13.45684 15.50683 13.83658 14.88543 141.9334   100
chol_method(X, w, y) 16.78841 17.30043 18.00251 17.69015 18.69025 21.1731   100

$`N=2000, P=500`
Unit: milliseconds
      expr      min       lq      mean     median        uq      max neval
inv_method(X, w, y) 4114.8367 5645.3391 5603.0128 5721.7032 5837.2093 7025.8754   100
lu_method(X, w, y)  492.1905 680.8708 699.6537 689.3546 713.6434 911.0215   100
chol_method(X, w, y) 627.7260 872.8408 870.3740 886.8396 910.2068 1052.0467   100
```

Figure 1 Performance benchmarking results

solution_1c.R

Shuchen

Wed Sep 13 00:57:14 2017

```
# SDS385 Exercise 1 Part C
# compare various matrix decomposition methods with inversion method
# and benchmarks performance of the Cholesky and LU versus inversion
```

```
library ( Matrix )
library ( microbenchmark ) # For benchmarking
```

```
# Inversion method function
inv_method <- function (X,W,y) {
  beta_hat <- solve (t(X) %*% W %*% X) %*% t(X) %*% W %*% y
  return (beta_hat)
}
```

```
# Cholesky decomposition
chol_method <- function (X,W,y) {
  A = t(X) %*% (X*diag (W)) # Efficient way of  $A = t(X) \%*\% W \%*\% X$  as  $W$  is diagonal
  b = t(X) %*% (y*diag (W)) # Avoid multiply by 0's

  # Cholesky decomposition of A
  U <- chol(A) # Upper Cholesky decomposition of A and  $U'U = A$ 

  # Replace  $Ax=b$  with  $U'Ux=b$ , solve  $U'z=b$  for z
  z <- solve(t(U))%*%b
  # Solve  $Ux = z$  for x ( $x=beta\_hat$ )
  beta_hat_chol <- solve(U) %*% z

  return(beta_hat_chol)
}
```

```
# LU method function
lu_method <- function (X,W,y) {
  A = t(X) %*% (X*diag (W)) # Efficient way of  $A = t(X) \%*\% W \%*\% X$  as  $W$  is diagonal
  b = t(X) %*% (y*diag (W)) # Avoid multiply by 0's

  # LU decomposition of A
  decomp <- lu(A)
  L <- expand ( decomp )$L # Upper triangular matrix
  U <- expand ( decomp )$U # Lower triangular matrix

  # Replace  $Ax=b$  with  $LUx=b$ , solve  $Lz=b$  for z
  z <- solve (L) %*% b
  # Solve  $Ux=z$  for x ( $x=beta\_hat$ )
  beta_hat_lu <- solve(U) %*% z

  return(beta_hat_lu)
}
```

```

# Simulate data from the linear model for a range of values of N and P
# Assume weights  $w_i$  are all 1, data are Gaussian

N <- c (20, 100, 400, 1200)
P <- N/4 #  $N \gg P$ 

res <- list () # Performance results
for (i in 1: length (N)) {
  n <- N[i]
  p <- P[i]
  print (n)

  # Set up matrices of size N, P parameters: (dummy data)
  X <- matrix ( rnorm (n*p), nrow =n, ncol =p)
  y <- rnorm (n)
  W <- diag (1, nrow =n)

  # Perform benchmarking:
  res[[i]] <- microbenchmark (
    inv_method (X,W,y),
    lu_method (X,W,y),
    chol_method (X,W,y),
    unit ='ms'
  )
}

```

```

## [1] 20
## [1] 100
## [1] 400
## [1] 1200

```

```
names(res) <- (c('N=20, P=5', 'N=100, P=25', 'N=400, P=100', 'N=1200, P=300'))
```

```
res # Display benchmarking results
```

```
## $`N=20, P=5`
## Unit: milliseconds
##           expr      min      lq      mean     median      uq
## inv_method(X, W, y) 0.076404 0.0798255 0.1334967 0.083532 0.089233
## lu_method(X, W, y) 0.163642 0.1687740 0.3182182 0.173336 0.197283
## chol_method(X, W, y) 0.157941 0.1625010 0.2971386 0.167633 0.188160
##      max neval
## 4.678906    100
## 12.778897    100
## 9.303074    100
##
## $`N=100, P=25`
## Unit: milliseconds
##           expr      min      lq      mean     median      uq
## inv_method(X, W, y) 0.628910 0.6788010 0.7098014 0.6899190 0.7058845
## lu_method(X, W, y) 0.269125 0.3030510 0.3616489 0.3204420 0.3572185
## chol_method(X, W, y) 0.317021 0.3549375 0.4032436 0.3680515 0.3894340
##      max neval
## 1.276065    100
## 2.023002    100
## 2.097696    100
##
## $`N=400, P=100`
## Unit: milliseconds
##           expr      min      lq      mean     median      uq
## inv_method(X, W, y) 31.444913 32.104898 34.481651 32.892888 33.919784
## lu_method(X, W, y) 4.365306 4.629015 5.098359 4.819455 5.217442
## chol_method(X, W, y) 5.685845 5.927887 6.412872 6.148548 6.615526
##      max neval
## 158.224091    100
## 8.047535    100
## 9.385750    100
##
## $`N=1200, P=300`
## Unit: milliseconds
##           expr      min      lq      mean     median      uq
## inv_method(X, W, y) 844.7061 986.6886 1122.2103 1196.6246 1209.0617
## lu_method(X, W, y) 106.2886 114.4154 145.2011 151.0524 153.7698
## chol_method(X, W, y) 134.2309 139.3554 176.9943 192.2784 195.7186
##      max neval
## 1344.5684    100
## 297.3477    100
## 337.5027    100
```

SDS385 Exercise 1 Part D ????????Still working on it

```
library(Matrix)
```

```
N <- 10000
```

```
P <- N/4
```

```
# simulate a sparse matrix X
```

```
X = matrix(rnorm(N*P), nrow=N)
```

```
mask = matrix(rbinom(N*P, 1, 0.05), nrow=N)
```

```
X = mask*X
```

```
X[1:10, 1:10] # quick visual check
```

```
##      [, 1] [, 2] [, 3]      [, 4] [, 5]      [, 6] [, 7]      [, 8] [, 9] [, 10]
## [1,]    0    0    0 3.0864389    0 0.000000    0 0.000000    0    0
## [2,]    0    0    0 0.0000000    0 0.000000    0 0.000000    0    0
## [3,]    0    0    0 0.0000000    0 0.000000    0 -0.9737562    0    0
## [4,]    0    0    0 0.8790703    0 0.000000    0 0.000000    0    0
## [5,]    0    0    0 0.0000000    0 0.000000    0 0.000000    0    0
## [6,]    0    0    0 0.0000000    0 0.000000    0 0.000000    0    0
## [7,]    0    0    0 0.0000000    0 0.000000    0 0.000000    0    0
## [8,]    0    0    0 0.0000000    0 1.892181    0 0.000000    0    0
## [9,]    0    0    0 0.0000000    0 0.000000    0 0.000000    0    0
## [10,]   0    0    0 0.0000000    0 0.000000    0 0.000000    0    0
```

```
# Or X <- Matrix(data=rnorm(N*P)*rbinom(N*P, 1, 0.01), nrow=N, sparse=TRUE)
```

Generalized Linear Regression

(A) The negative log likelihood is

$$\begin{aligned}
 l(\beta) &= -\log\left\{\prod_{i=1}^N p(y_i | \beta)\right\} = -\log\left\{\prod_{i=1}^N \binom{m_i}{y_i} w_i^{y_i} (1 - w_i)^{m_i - y_i}\right\} \\
 &= -\sum_{i=1}^N \log\left\{\binom{m_i}{y_i} w_i^{y_i} (1 - w_i)^{m_i - y_i}\right\} \\
 &= -\sum_{i=1}^N \left\{\log\binom{m_i}{y_i} + y_i \log(w_i) + (m_i - y_i) \log(1 - w_i)\right\}
 \end{aligned}$$

$$\text{With } w_i = \frac{1}{1 + \exp(-x_i^T \beta)} \Rightarrow \exp(-x_i^T \beta) = \frac{1}{w_i} - 1$$

Taking derivative of $l(\beta)$ w.r.t β :

$$\nabla l(\beta) = -\sum_{i=1}^N \left(\frac{y_i}{w_i} \frac{\partial w_i}{\partial \beta} - \frac{m_i - y_i}{1 - w_i} \frac{\partial w_i}{\partial \beta} \right)$$

$$\text{Since } \frac{\partial w_i}{\partial \beta} = -(1 + \exp(-x_i^T \beta))^{-2} \frac{\partial}{\partial \beta} \exp(-x_i^T \beta)$$

$$= \frac{-\exp(-x_i^T \beta)(-x_i)}{(1 + \exp(-x_i^T \beta))^2} = \frac{x_i \exp(-x_i^T \beta)}{(1 + \exp(-x_i^T \beta))^2} = w_i^2 x_i \left(\frac{1}{w_i} - 1 \right) = w_i x_i (1 - w_i)$$

$$\Rightarrow \nabla l(\beta) = -\sum_{i=1}^N \left(\frac{y_i}{w_i} w_i x_i (1 - w_i) - \frac{m_i - y_i}{1 - w_i} w_i x_i (1 - w_i) \right)$$

$$= -\sum_{i=1}^N (y_i x_i (1 - w_i) - (m_i - y_i) w_i x_i) = -\sum_{i=1}^N (y_i - m_i w_i) x_i$$

In matrix form

$$= -X^T (y - mw) = X^T (mw - y)$$